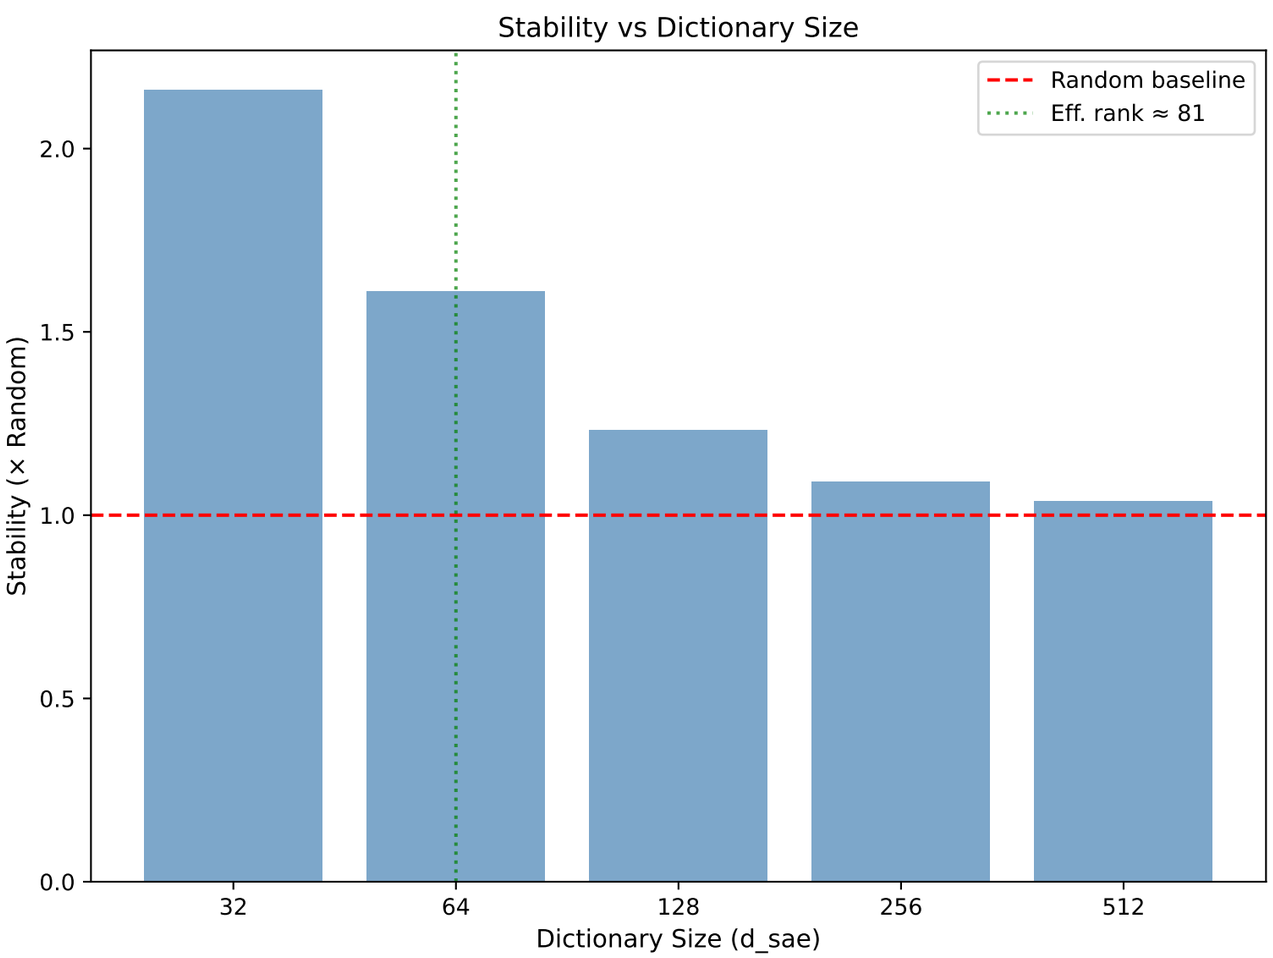

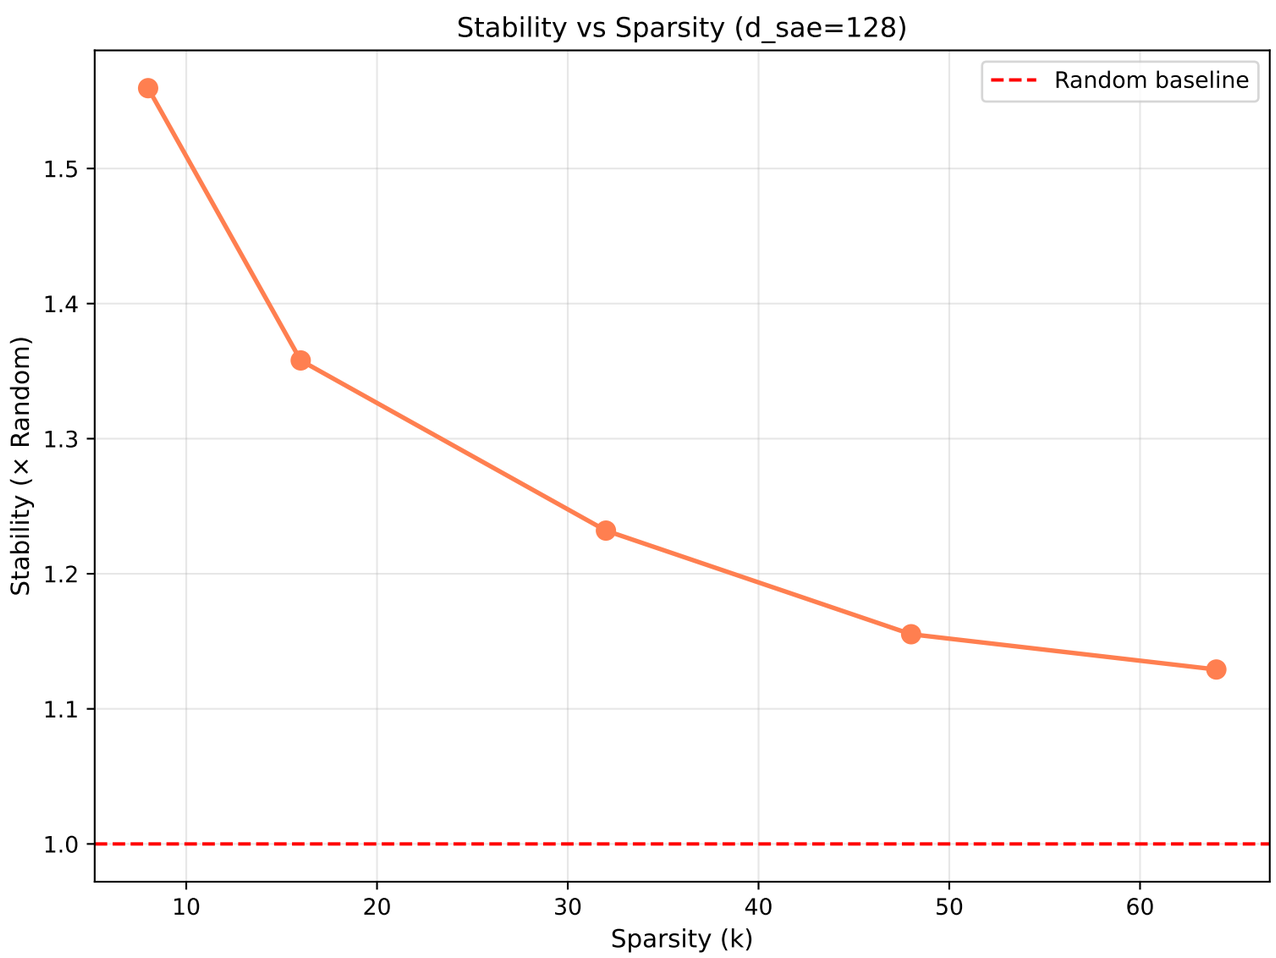

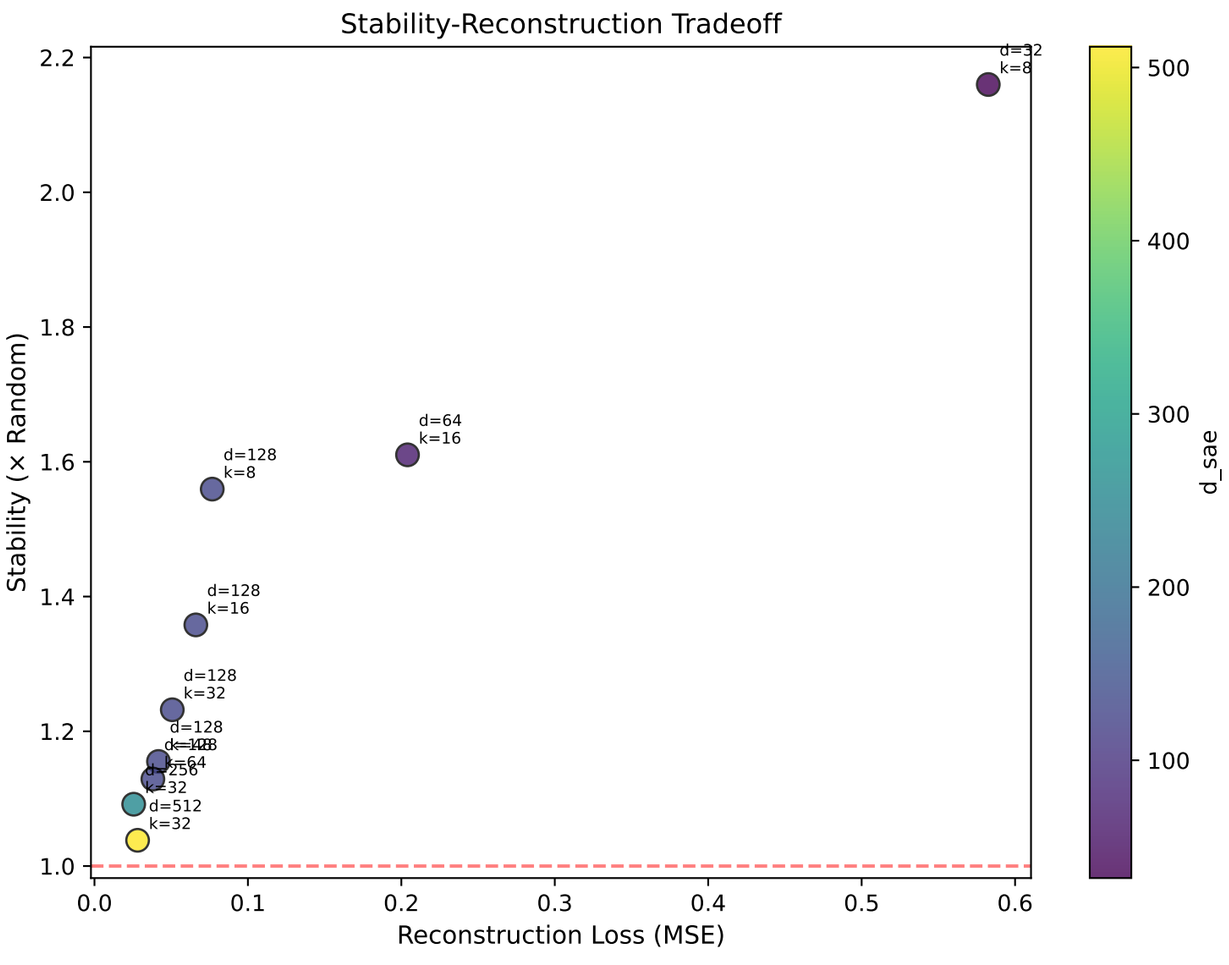

KEY FINDINGS

1. STABILITY-RECONSTRUCTION TRADEOFF
   - Smaller SAEs → Higher stability, worse reconstruction
   - Larger SAEs → Lower stability (≈ random), better reconstruction
   - Matched regime (d_sae ≈ eff_rank) offers best balance

2. STABILITY DECREASES WITH SPARSITY (k)
   - Lower k → Higher stability (more constrained)
   - Higher k → Lower stability (more freedom)
   - This is OPPOSITE to LLM findings!

3. FEATURE-LEVEL STABILITY IS UNIFORM
   - No predictor (frequency, magnitude, task correlation) significantly predicts feature stability
   - Stability is a GLOBAL property, not feature-specific

4. TASK-DEPENDENT STABILITY
   - On algorithmic tasks: constraint = stability
   - On LLMs: may have optimal sparsity for "correct" features
   - Semantic structure may be required for non-monotonic stability

IMPLICATIONS

- SAE stability findings from LLMs may NOT transfer to algorithmic tasks
- For interpretability: use matched regime (d_sae ≈ eff_rank)
- Stability is fundamentally about CONSTRAINT, not correctness